

Homework 2, STATS 315A

Stanford University, Winter 2019

Joe Higgins

Question 2

This question will have you write an R function to efficiently perform forward stepwise linear regression. It requires that you understand the material on pages 13–15 of the chapter 3 lecture notes. The setup is linear regression with an $N \times p$ matrix X and a response vector y . You always include an intercept in your models, which is not included in X .

- (a) Let \tilde{X} be the matrix X with each of the columns mean centered. What is the fitted intercept in the regression of y on \tilde{X} ? How do the coefficients of X in the regression of y on X compare with the coefficients of \tilde{X} ?

#stuff goes here

Question 7

7. Obtain the zipcode train and test data from the ESL website.
- Compare the test performance of a) linear regression b) linear discriminant analysis and c) multiclass linear logistic regression.
 - For a) and c), use the package glmnet (available in R, matlab and python) to run elastic-net regularized versions of each (use $\alpha = 0.3$). For these two, plot the test error as a functions of the training R2 for a) and D2 for c) (% training deviance explained).
 - In ii., what is the optimization problem being solved?

```
rm(list = ls())

read_data_as_matrix <- function(file_string){
  table <- read.table(file_string)
  output <- matrix(, nrow=dim(table)[1], ncol=dim(table)[2])
  for(i in 1:ncol(table)){
    output[,i] <- table[,i]
  }
  return(output)
}

get_accuracy <- function(y_hat, y){
  correct <- y_hat == y
  pct_correct <- sum(correct)/length(correct)
  return(pct_correct)
}

train <- read_data_as_matrix("zip.train")
y_train <- train[,1]
x_train <- train[,2:dim(train)[2]]
test <- read_data_as_matrix("zip.test")
y_test <- test[,1]
x_test <- test[,2:dim(test)[2]]
```

```

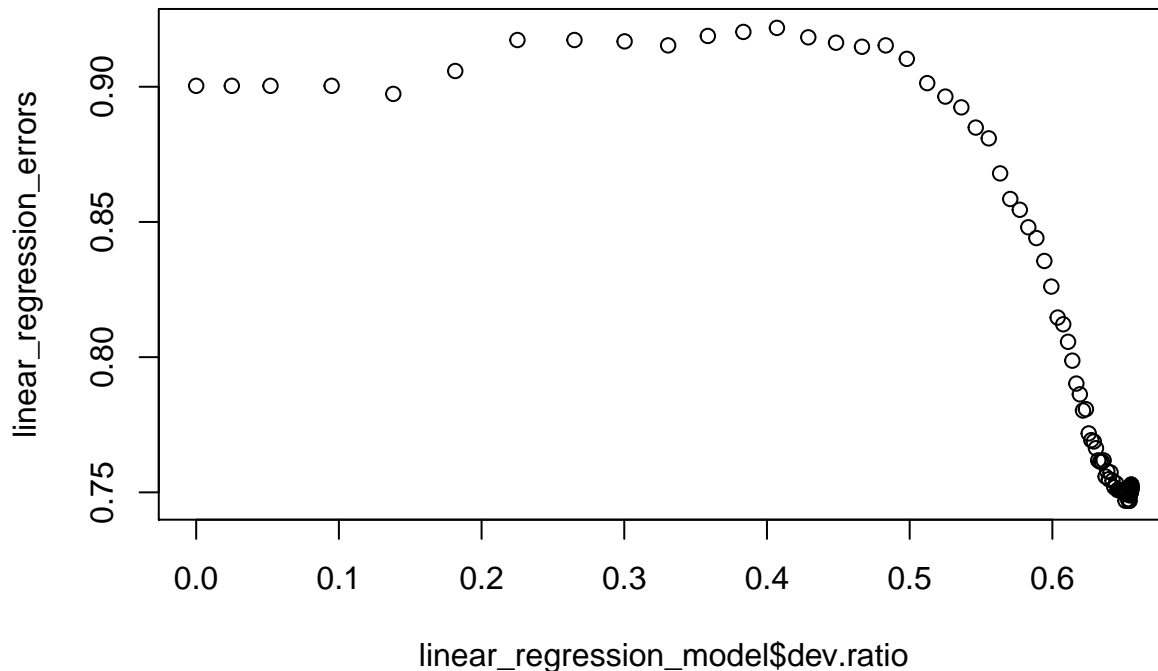
linear_regression_model <- glmnet(x_train, y_train, family=c("gaussian"), alpha = 0.3)
lda_model <- lda(y_train ~ ., data=data.frame(x_train), na.action="na.omit", CV=FALSE)
multinomial_regression_model <- glmnet(x_train, y_train, family=c("multinomial"), alpha = 0.3)

linear_regression_output <- predict(linear_regression_model, x_test, type=c("response"))
linear_regression_prdct <- round(linear_regression_output)
linear_regression_prdct[linear_regression_prdct == -1] <- 0
multinomial_regression_output <- predict(multinomial_regression_model, x_test, type=c("response"))
multinomial_regression_prdct <- apply(
  multinomial_regression_output, 3, function(x) apply(
    x, 1, function(y) which.max(y) - 1
  )
)
lda_output <- predict(lda_model, data.frame(x_test))
lda_prdct <- lda_output$class

linear_regression_errors <- apply(linear_regression_prdct, 2, function(x) 1 - get_accuracy(x, y_test))
multinomial_regression_errors <- apply(multinomial_regression_prdct, 2, function(x) 1 - get_accuracy(x,
lda_accuracy <- get_accuracy(lda_prdct, y_test)

#is dev ratio the R2 for linear regression?
#what is the optimization being done...
plot(linear_regression_model$dev.ratio, linear_regression_errors)

```



```

plot(multinomial_regression_model$dev.ratio, multinomial_regression_errors)

```

